

Group Work 05

```
library(bayesrules) # R package for our textbook
library(tidyverse) # Collection of packages for tidying and plotting data
library(janitor) # Helper functions like tabyl
library(rstan) # for fitting models
library(rstanarm) # for fitting standard regression models
library(broom.mixed) # for tidy() function
library(bayesplot) # helpful plotting functions
library(tidybayes) # helpful for wrangling Bayesian model output
```

Note

I've given you most of the R code you need for today to help you get through the questions more quickly and focus on the concepts. Make sure you understand what all the code is doing, enough so that you can replicate it on your own in the future.

In code chunks where you need to fill in the blank, remember to remove `#! eval: false` to get your document to render correctly

Throughout these problems, you'll be modeling the daily 3pm temperature in Perth, Australia, based on the 9am temperature. We'll use the `weather_perth` data in the `{bayesrules}` package

```
data("weather_perth")
```

We'll use the following Bayesian regression model:

$$Y_i | \beta_0, \beta_1, \sigma \sim N(\mu_i, \sigma^2)$$

$$\mu_i = \beta_0 + \beta_1 X_i$$

$$\beta_{0c} \sim N(30, 7^2)$$

$$\beta_1 \sim N(2, 1^2)$$

$$\sigma \sim \text{Exp}(0.1)$$

1 Prior understanding

- (a) Start by summarizing our prior knowledge of $\beta_{0c}, \beta_1, \sigma$. Fill in the blanks to plot the appropriate prior distributions, then provide a 1 sentence summary of each parameter.

```
p1 <- plot_normal(____, ____ ) +
  labs(x = expression(beta[0c]), y = "density")

p2 <- plot_normal(____, ____ ) +
  labs(x = expression(beta[1]), y = "density")

p3 <- plot_gamma(____, ____ ) +
  labs(x = expression(sigma), y = "density")

p1 + p2 + p3
```

- (b) The prior models help us understand each individual parameter, $\beta_{0c}, \beta_1, \sigma$. To better understand how these work together in the model of 3pm temperatures, we can simulate a sample of prior plausible parameter sets using the `stan_glm()` function in the `{rstanarm}` package. Fill in the blanks in the code, then annotate each line of code with a comment about what it does.

```
# Simulate the prior model
temp_prior <- stan_glm(
  temp3pm ~ temp9am, # your comment here
  data = weather_perth,
  family = gaussian, # your comment here
  prior_intercept = normal(____, ____), # your comment here
  prior = normal(____, ____), # your comment here
  prior_aux = exponential(____), # your comment here
  prior_PD = TRUE, # your comment here
  chains = 4, iter = 5000*2, seed = 84735) # your comment here
```

- (c) `temp_prior` should contain 20,000 plausible parameter sets. It's important to check our priors. The following chunk simulates 200 prior plausible regression lines $\beta_0 + \beta_1 X$, 1 from each of 200 prior plausible parameter sets in `temp_prior`. Run it a few times and summarize your observations.

```
# 200 prior model lines
weather_perth %>%
  add_epred_draws(temp_prior, ndraws = 200) %>%
  ggplot(aes(x = temp9am, y = temp3pm)) +
    geom_line(aes(y = .epred, group = .draw), alpha = 0.15)
```

- (d) The plot in part c illuminates prior plausible regression lines $\beta_0 + \beta_1 X$. To explore our prior understanding of potential variability from this line, σ , the following chunk simulates 4 prior plausible sets of temperature data, 1 from each of 4 prior plausible parameter sets in temp_prior. Run it a few times and summarize your observations.

```
# Plot 4 prior simulated datasets
weather_perth %>%
  add_predicted_draws(temp_prior, ndraws = 4) %>%
  ggplot(aes(x = temp9am, y = temp3pm)) +
    geom_point(aes(y = .prediction, group = .draw), size = 0.1) +
    facet_wrap(~ .draw)
```

- (e) Explain the difference between the analysis in (c) and in (d)

2 Simulate the posterior

- (a) Next, let's simulate the posterior. We don't have to start from scratch. Rather, we can `update()` our prior simulation using `prior_PD = FALSE`.

```
# Simulate the posterior
```

- (b) To ensure that we can “trust” the MCMC posterior simulation results, create trace plots and density plots of the multiple chains.

```
# your code here
```

- (c) Just as we did for the prior, we can better understand our posterior model by simulating 50 posterior plausible regression lines $\beta_0 + \beta_1 X$. Run this chunk a few times. Summarize your observations as well as how our understanding evolved from the prior to the posterior.

```
## 50 simulated posterior model lines

___ %>%
  ___(___, ndraws = 50) %>%
  ggplot(aes(x = ___, y = ___)) +
    geom_line(aes(y = ___, group = ___), alpha = 0.15) +
    geom_point(data = weather_perth, size = 0.05)
```

- (d) Check out a posterior summary table. Do we have ample evidence of a positive association between 3pm and 9am temperatures, i.e. that the warmer it is in the morning, the warmer it tends to be in the afternoon? Explain.

```
# Get posterior summaries of the "fixed" parameters (betas)
# and "auxiliary" parameter (sigma)
tidy(temp_posterior, effects = c("fixed", "aux"),
     conf.int = TRUE, conf.level = 0.80)
```

3 Posterior prediction

- (a) Suppose it's 20 degrees at 9am. Simulate, plot, and describe a posterior predictive model for the 3pm temperature. HINTS: Predict a 3pm temperature from each of the 20,000 parameter sets. To work with `(Intercept)`, put a backtick at the start and end.

```
# Store the MCMC chains as a data frame
temp_chains <- as.data.frame(temp_posterior)

# Simulate the posterior predictive model
set.seed(84735)
predict_20 <- ___ %>%
  ___(y_prediction = ___)

# Plot the posterior predictive model
ggplot(predict_20, aes(x = ___)) +
  geom_density()
```

- (b) It's important to understand where the predictions come from. BUT there's also a shortcut. Use the `posterior_predict()` function to simulate the posterior predictive model of 3pm temperature when it's 20 degrees at 9am. Further, use `posterior_interval()` and `mcmc_dens()` to summarize the posterior (use a 90% level). Examine the results and confirm that they're similar to your "from scratch" simulation.

```
# Simulate a set of predictions
set.seed(84735)
shortcut_prediction <- posterior_predict(
  temp_posterior, newdata = data.frame(temp9am = 20))

# Construct a 90% posterior credible interval
posterior_interval(shortcut_prediction, prob = ___)
```

```
# Plot the approximate predictive model
mcmc_areas(shortcut_prediction, prob = ___) +
  scale_x_continuous(breaks = seq(15, 35, by = 5)) +
  xlab("3pm temperature")
```

4 Posterior predictive check

In the question above, we considered posterior predictions for a single (new) data point. We can also use posterior predictions *on the entire original dataset* to evaluate how well our model is performing.

- (a) The code below (1) pulls out the first set of $(\beta_0, \beta_1, \sigma)$ from our MCMC sample and then (2) simulates a prediction \hat{y} based on those estimates, for every point in the original dataset. Fill in the blanks and then take a peek at `one_simulation` to make sure you understand what it contains.

```
first_set <- head(temp_chains, 1)

beta_0 <- first_set$`Intercept`
beta_1 <- first_set$temp9am
sigma <- first_set$sigma

set.seed(84735)
one_simulation <- weather_perth %>%
  mutate(mu = ___ + ___ * ___,
          simulated_temp3am = rnorm(500, mean = mu, sd = ___)) %>%
  select(temp9am, temp3pm, simulated_temp3pm)
```

Next, run this chunk of code to create a density plot of the predictions compared to the observed.

```
ggplot(one_simulation, aes(x = temp9am)) +
  geom_density(color = "lightblue") +
  geom_density(aes(x = rides), color = "darkblue")
```

- (b) It's *really important* to have a solid understanding of what a posterior predictive check is doing under the hood. But one of the benefits of using `{rstanarm}` is that, since all of the models have a consistent structure, we can make use of shortcut functions that exist in other packages. Use `pp_check` below to examine 50 simulated samples. Write a 1-sentence summary of your findings.

```
# Examine 50 of the 20000 simulated samples  
pp_check(temp_model, nreps = 50)
```